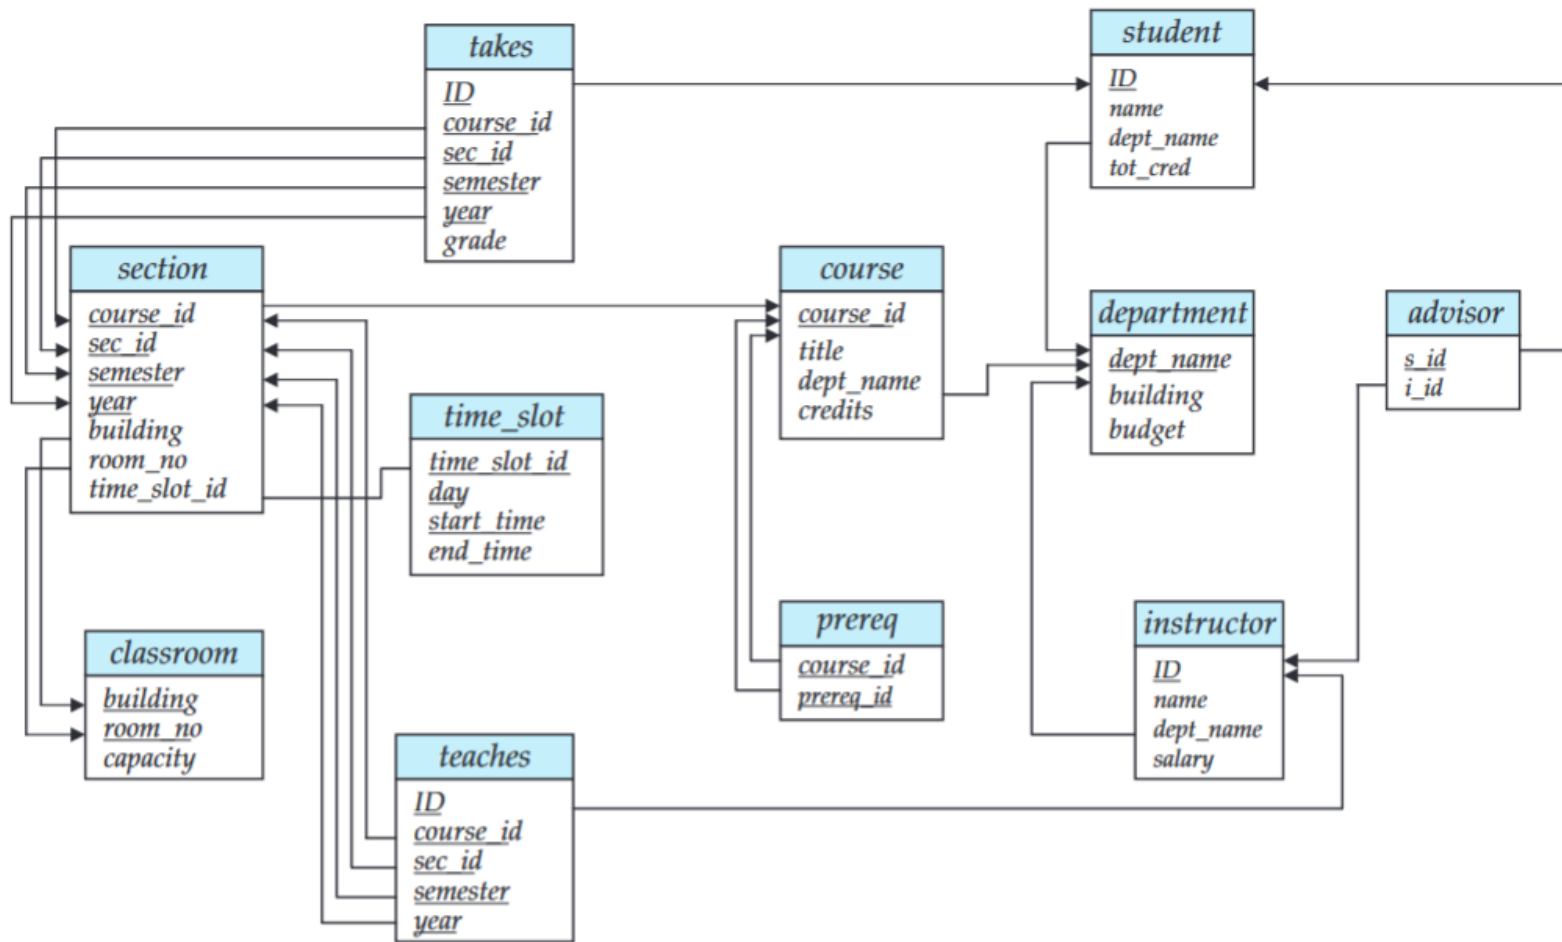


# SQL



**Figure 2.8** Schema diagram for the university database.

- 3.1 Write the following queries in SQL, using the university schema. (We suggest you actually run these queries on a database, using the sample data that we provide on the Web site of the book, [db-book.com](http://db-book.com). Instructions for setting up a database, and loading sample data, are provided on the above Web site.)
- Find the titles of courses in the Comp. Sci. department that have 3 credits.
  - Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.
  - Find the highest salary of any instructor.
  - Find all instructors earning the highest salary (there may be more than one with the same salary).
  - Find the enrollment of each section that was offered in Autumn 2009.
  - Find the maximum enrollment, across all sections, in Autumn 2009.
  - Find the sections that had the maximum enrollment in Autumn 2009.

- a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

- a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

```
select    title  
from      course  
where     dept_name = 'Comp. Sci.'  
           and credits = 3
```

- b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.

- b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.  
This query can be answered in several different ways. One way is as follows.

```
select    distinct student.ID
from      (student join takes using(ID))
            join (instructor join teaches using(ID))
            using(course_id, sec_id, semester, year)
where      instructor.name = 'Einstein'
```

As an alternative to the **join .. using** syntax above the query can be written by enumerating relations in the **from** clause, and adding the corresponding join predicates on *ID*, *course\_id*, *section\_id*, *semester*, and *year* to the **where** clause.

Note that using natural join in place of **join .. using** would result in equating student *ID* with instructor *ID*, which is incorrect.

c. Find the highest salary of any instructor.



- c. Find the highest salary of any instructor.

```
select max(salary)  
from instructor
```

- d. Find all instructors earning the highest salary (there may be more than one with the same salary).

- d. Find all instructors earning the highest salary (there may be more than one with the same salary).

```
select    ID, name  
from      instructor  
where     salary = (select max(salary) from instructor)
```

- e. Find the enrollment of each section that was offered in Autumn 2009.

- e. Find the enrollment of each section that was offered in Autumn 2009. One way of writing the query is as follows.

```
select    course_id, sec_id, count(ID)
from      section natural join takes
where     semester = 'Autumn'
and       year = 2009
group by course_id, sec_id
```

Note that if a section does not have any students taking it, it would not appear in the result. One way of ensuring such a section appears with a count of 0 is to replace **natural join** by the **natural left outer join** operation, covered later in Chapter 4. Another way is to use a subquery in the **select** clause, as follows.

- e. Find the enrollment of each section that was offered in Autumn 2009.  
One way of writing the query is as follows.

```
select   course_id, sec_id,
          (select count(ID)
from    takes
where   takes.year = section.year
          and takes.semester = section.semester
          and takes.course_id = section.course_id
          and takes.section_id = section.section_id)
from    section
where   semester = 'Autumn'
and     year = 2009
```

Note that if the result of the subquery is empty, the aggregate function **count** returns a value of 0.

f. Find the maximum enrollment, across all sections, in Autumn 2009.

- f. Find the maximum enrollment, across all sections, in Autumn 2009.  
One way of writing this query is as follows:

```
select max(enrollment)  
from (select   count(ID) as enrollment  
       from     section natural join takes  
       where    semester = 'Autumn'  
       and      year = 2009  
       group by course_id, sec_id)
```



- g. Find the sections that had the maximum enrollment in Autumn 2009.

- g. Find the sections that had the maximum enrollment in Autumn 2009. The following answer uses a **with** clause to create a temporary view, simplifying the query.

```
with sec_enrollment as (  
    select    course_id, sec_id, count(ID) as enrollment  
    from      section natural join takes  
    where     semester = 'Autumn'  
    and       year = 2009  
    group by course_id, sec_id)  
select  course_id, sec_id  
from    sec_enrollment  
where   enrollment = (select max(enrollment) from sec_enrollment)
```

It is also possible to write the query without the **with** clause, but the subquery to find enrollment would get repeated twice in the query.

3.4 Write the following inserts, deletes or updates in SQL, using the university schema.

- a. Increase the salary of each instructor in the Comp. Sci. department by 10%.
- b. Delete all courses that have never been offered (that is, do not occur in the *section* relation).
- c. Insert every student whose *tot\_cred* attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.

- a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

- a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

```
update instructor  
set      salary = salary * 1.10  
where dept_name = 'Comp. Sci.'
```

- b. Delete all courses that have never been offered (that is, do not occur in the *section* relation).

- b. Delete all courses that have never been offered (that is, do not occur in the *section* relation).

```
delete from course  
where course_id not in  
      (select course_id from section)
```

- c. Insert every student whose *tot\_cred* attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.



- c. Insert every student whose *tot\_cred* attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.

```
insert into instructor  
select ID, name, dept_name, 10000  
from student  
where tot_cred > 100
```

- 3.6 Suppose that we have a relation *marks*(ID, *score*) and we wish to assign grades to students based on the score as follows: grade *F* if  $score < 40$ , grade *C* if  $40 \leq score < 60$ , grade *B* if  $60 \leq score < 80$ , and grade *A* if  $80 \leq score$ . Write SQL queries to do the following:
- Display the grade for each student, based on the *marks* relation.

- a. Display the grade for each student, based on the *marks* relation.

```
select ID,  
       case  
         when score < 40 then 'F'  
         when score < 60 then 'C'  
         when score < 80 then 'B'  
         else 'A'  
       end  
from marks
```

b. Find the number of students with each grade.

b. Find the number of students with each grade.

```
with      grades as  
(  
  select   ID,  
            case  
              when score < 40 then 'F'  
              when score < 60 then 'C'  
              when score < 80 then 'B'  
              else 'A'  
            end as grade  
  from marks  
)  
select    grade, count(ID)  
from      grades  
group by grade
```

As an alternative, the **with** clause can be removed, and instead the definition of *grades* can be made a subquery of the main query.

3.7 The SQL **like** operator is case sensitive, but the `lower()` function on strings can be used to perform case insensitive matching. To show how, write a query that finds departments whose names contain the string “sci” as a substring, regardless of the case.

```
select  dept_name  
from    department  
where   lower(dept_name) like '%sci%'
```

**3.12** Write the following queries in SQL, using the university schema.

- a. Create a new course “CS-001”, titled “Weekly Seminar”, with 0 credits.
- b. Create a section of this course in Autumn 2009, with *section.id* of 1.
- c. Enroll every student in the Comp. Sci. department in the above section.
- d. Delete enrollments in the above section where the student’s name is Chavez.
- e. Delete the course CS-001. What will happen if you run this delete statement without first deleting offerings (sections) of this course.
- f. Delete all *takes* tuples corresponding to any section of any course with the word “database” as a part of the title; ignore case when matching the word with the title.



a. SQL query:

```
insert into course  
values ('CS-001', 'Weekly Seminar', 'Comp. Sci.', 0)
```

b. SQL query:

```
insert into section  
values ('CS-001', 1, 'Autumn', 2009, null, null, null)
```

c. SQL query:

```
insert into takes  
  select id, 'CS-001', 1, 'Autumn', 2009, null  
  from student  
  where dept_name = 'Comp. Sci.'
```

d. SQL query:

```
delete from takes
where course_id = 'CS-001' and section_id = 1 and
      year = 2009 and semester = 'Autumn' and
      id in (select id
             from student
             where name = 'Chavez')
```

Note that if there is more than one student named Chavez, all such students would have their enrollments deleted. If we had used = instead of **in**, an error would have resulted if there were more than one student named Chavez.

e. SQL query:

```
delete from takes  
where course_id = 'CS-001'
```

```
delete from section  
where course_id = 'CS-001'
```

```
delete from course  
where course_id = 'CS-001'
```

If we try to delete the course directly, there will be a foreign key violation because *section* has a foreign key reference to *course*; similarly, we have to delete corresponding tuples from *takes* before deleting sections, since there is a foreign key reference from *takes* to *section*. As a result of the foreign key violation, the transaction that performs the delete would be rolled back.

f. SQL query:

```
delete from takes  
where course_id in  
    (select course_id  
     from course  
     where lower(title) like '%database%')
```

# Employee database

*employee* (*employee\_name*, *street*, *city*)  
*works* (*employee\_name*, *company\_name*, *salary*)  
*company* (*company\_name*, *city*)  
*manages* (*employee\_name*, *manager\_name*)

Give an SQL schema definition for the employee database. Choose an appropriate domain for each attribute and an appropriate primary key for each relation schema

<b>create table</b>	<i>employee</i>
<i>(employee_name</i>	<b>varchar(20),</b>
<i>street</i>	<b>char(30),</b>
<i>city</i>	<b>varchar(20),</b>
<b>primary key</b>	<i>(employee_name))</i>



<b>create table</b>	<i>works</i>
<i>(employee_name</i>	<i>person_names,</i>
<i>company_name</i>	<b>varchar</b> (20),
<i>salary</i>	<b>numeric</b> (8, 2),
<b>primary key</b>	<i>(employee_name))</i>

<b>create table</b>	<i>company</i>
<i>(company_name</i>	<b>varchar</b> (20),
<i>city</i>	<b>varchar</b> (20),
<b>primary key</b>	<i>(company_name))</i>

```
create table      manages  
(employee_name   varchar(20),  
  manager_name  varchar(20),  
  primary key   (employee_name))
```

Give an expression in SQL for each of the following queries.

- a. Give all employees of First Bank Corporation a 10 percent raise.
- b. Give all managers of First Bank Corporation a 10 percent raise.
- c. Delete all tuples in the *works* relation for employees of Small Bank Corporation.

Answer:

- a. Give all employees of First Bank Corporation a 10-percent raise. (the solution assumes that each person works for at most one company.)

```
update works  
set salary = salary * 1.1  
where company_name = 'First Bank Corporation'
```

b. Give all managers of First Bank Corporation a 10-percent raise

```
update works  
set salary = salary * 1.1  
where employee_name in (select manager_name  
                           from manages)  
      and company_name = 'First Bank Corporation'
```

- c. Delete all tuples in the *works* relation for employees of Small Bank Corporation.

```
delete from works  
where company_name = 'Small Bank Corporation'
```

# Library database

*member(memb\_no, name, age)*  
*book(isbn, title, authors, publisher)*  
*borrowed(memb\_no, isbn, date)*

Consider the library database. Write the following queries in SQL.

- a. Print the names of members who have borrowed any book published by “McGraw-Hill”.
- b. Print the names of members who have borrowed all books published by “McGraw-Hill”.
- c. For each publisher, print the names of members who have borrowed more than five books of that publisher.
- d. Print the average number of books borrowed per member. Take into account that if an member does not borrow any books, then that member does not appear in the *borrowed* relation at all.



- a. Print the names of members who have borrowed any book published by McGraw-Hill.

```
select name
from member m, book b, borrowed l
where m.memb_no = l.memb_no
       and l.isbn = b.isbn and
          b.publisher = 'McGrawHill'
```

- b. Print the names of members who have borrowed all books published by McGraw-Hill. (We assume that all books above refers to all books in the *book* relation.)

```
select distinct m.name  
from member m  
where not exists  
    ((select isbn  
    from book  
    where publisher = 'McGrawHill')  
    except  
    (select isbn  
    from borrowed l  
    where l.memb_no = m.memb_no))
```

- c. For each publisher, print the names of members who have borrowed more than five books of that publisher.

```
select publisher, name
from (select publisher, name, count (isbn)
      from member m, book b, borrowed l
      where m.memb_no = l.memb_no
      and l.isbn = b.isbn
      group by publisher, name) as
      membpub(publisher, name, count_books)
where count_books > 5
```

d. Print the average number of books borrowed per member.

```
with memcount as  
    (select count(*)  
     from member)  
select count(*)/memcount  
from borrowed
```

Note that the above query ensures that members who have not borrowed any books are also counted. If we instead used **count(distinct *memb\_no*)** from *borrowed*, we would not account for such members.

### 3.22 Rewrite the **where** clause

**where unique** (**select** *title* **from** *course*)

without using the **unique** construct.

**Answer:**

```
where  
  (select count(title)  
   from course) =  
  (select count (distinct title)  
   from course))
```

### 3.24 Consider the query:

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
     from instructor  
     group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
     from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value >= dept_total_avg.value;
```

Rewrite this query without using the **with** construct.

**Answer:**

There are several ways to write this query. One way is to use subqueries in the where clause, with one of the subqueries having a second level subquery in the from clause as below.

```

select distinct dept_name d
from instructor i
where
    (select sum(salary)
     from instructor
     where department = d)
    >=
    (select avg(s)
     from
        (select sum(salary) as s
         from instructor
         group by department))

```

Note that the original query did not use the *department* relation, and any department with no instructors would not appear in the query result. If we had written the above query using *department* in the outer **from** clause, a department without any instructors could appear in the result if the condition were  $\leq$  instead of  $\geq$ , which would not be possible in the original query.

As an alternative, the two subqueries in the where clause could be moved into the from clause, and a join condition (using  $\geq$ ) added.