

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern and dynamic visual effect.

Introduction to Query Processing and Query Optimization

Outline

- ▶ Overview
- ▶ Measures of Query Cost
- ▶ Query Optimization

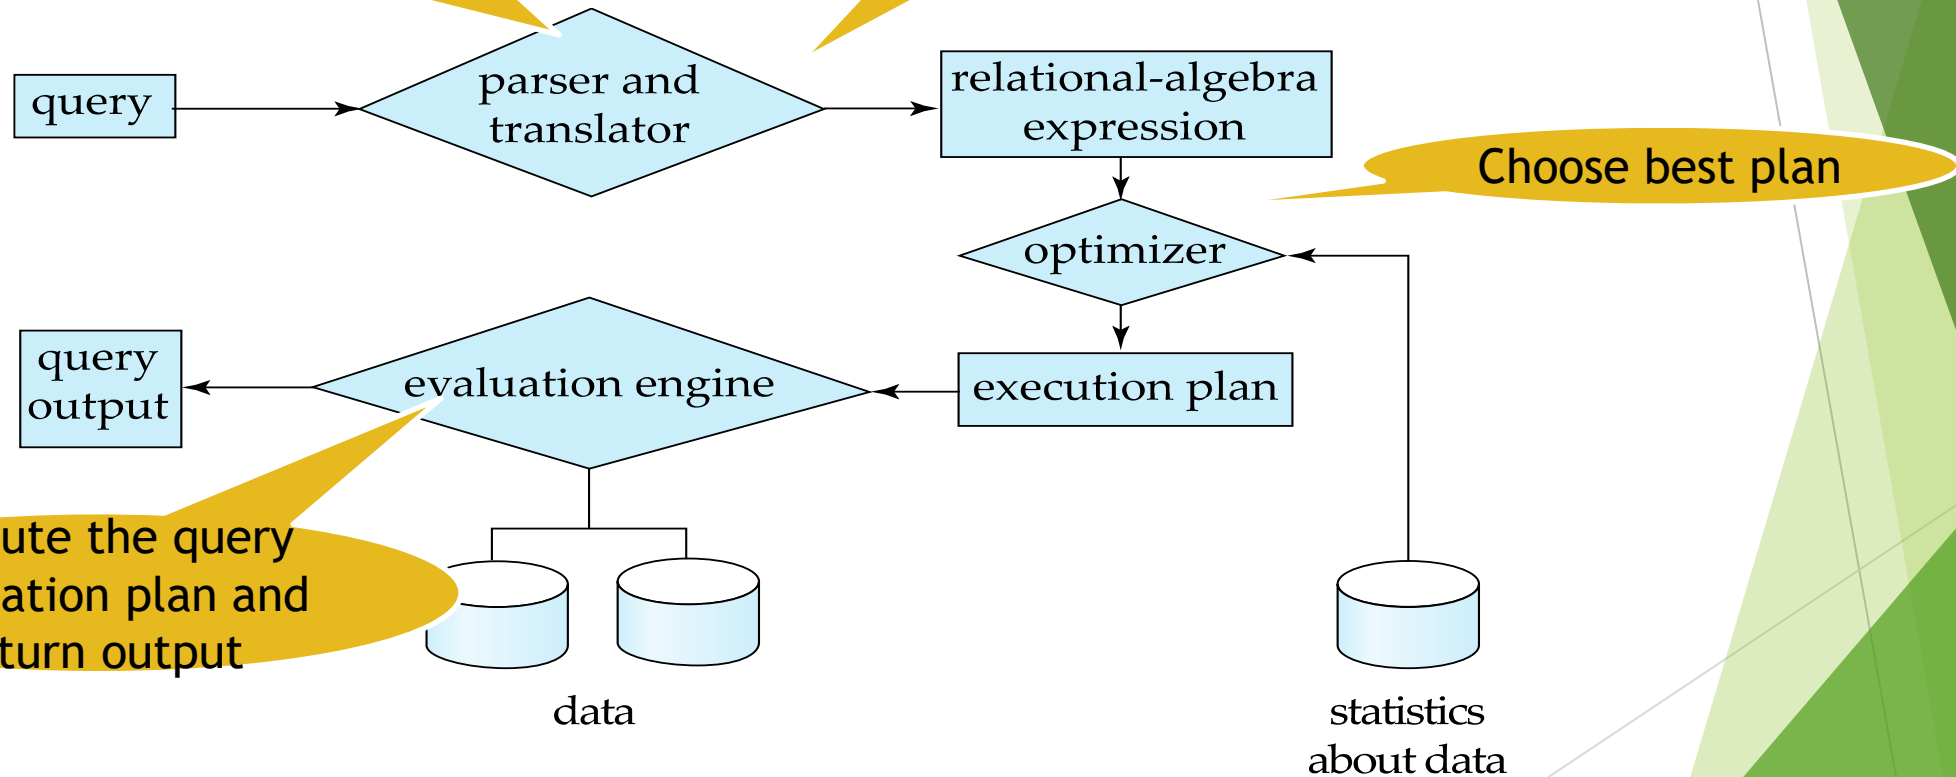
What is Query Processing?

- Query processing: Activities involved in extracting data from a database.
- Three basic steps:
 1. Parsing and Translation
 2. Optimization
 3. Evaluation

Steps in Query Processing

Parser check the syntax of query and verifies attribute name and relationship name

Translator translate the query into its internal form



Execute the query evaluation plan and return output

statistics about data

Measures of Query Cost

- Cost is generally measured as total elapsed time for answering query
- Many factors contribute to time cost
 1. Disk accesses (Time to process a data request and retrieve data from the storage device)
 2. CPU (time to execute a query)
 3. Network communication cost
- Disk access is the predominant cost, and is also relatively easy to estimate.
- Cost to write a block is greater than cost to read a block
 - data is read back after being written to ensure that the write was successful

Measures of Query Cost

- ▶ For simplicity we just use the number of block transfers from disk and the number of seeks as the cost measures

tT – time to transfer one block

tS – time for one seek

- ▶ Cost for b block transfers plus S seeks

$$b * tT + S * tS$$

- ▶ We ignore CPU costs for simplicity

Real systems do take CPU cost into account

We do not include cost to writing output to disk in our cost

Select operation

- **Symbol:** σ
- **Notation:** σ condition (Relation)
- **Operation :** Select tuple from a relation that satisfy a given condition.

- Search algorithm
 1. Linear search (A1)
 2. Binary search (A2)

Linear search (A1)

- ▶ It Scan each file block and test all records to see whether they satisfy the selection condition.
 - ▶ Cost estimate = br block transfers
 - br denotes number of blocks containing records from relation r
 - ▶ If selection is on a key attribute (primary key), then system can stop on finding record
 - cost = $(br / 2)$ block transfers
 - ▶ Linear search can be applied regardless of
 - Selection condition or
 - Ordering of records in the file, or
 - Availability of indices
- ▶ This algorithm is slower than binary search algorithm.

Binary search (A2)

- ▶ Is used when selection is an **equality comparison** on the **primary key** and **relation is sorted on primary key attribute**.
- ▶ Cost of binary search = $\lceil \log_2(br) \rceil$
br denotes number of blocks containing records from relation r
- ▶ If the selection is on **non primary attribute** then **multiple block may contains required records** , then the **cost of scanning** such block need to be **added** to the cost estimate.
- ▶ This algorithm is faster than linear search algorithm

Evaluation of expressions

➤ Method

1. Materialization
2. Pipelining

Materialization

- **Materialized evaluation:** evaluate one operation at a time, starting at the lowest-level. (from bottom and perform the inner most operations first)
- The intermediate results of each operation is materialized (store in temporary relation) and become input for subsequent(evaluate next-level operations).
- The cost of materialization is the sum of the individual operations plus the cost of writing the intermediate results to disk.
- ▶ The problem ;
 1. Creates lots of temporary relation
 2. Perform lots of I/O operation

Pipelining

- ▶ It evaluate several operations simultaneously, passing the results of one operation on to the next.
- ▶ To reduce number of intermediate temporary relations , we pass results of one operation to the next operation in the pipeline.
- ▶ Combining operations into a pipeline eliminates the cost of reading and writing temporary relations.
- ▶ Much cheaper than materialization: no need to store a temporary relation to disk.
- ▶ Pipelines can be executed in two ways:
- ▶ **Demand driven** –system makes request for tuples from the operation at the top of pipeline
- ▶ **Producer driven** – Operation do not wait for request to produce tuple but generate the tuples eagerly.

Query Optimization

Process of selecting the most efficient query evaluation plan

Query Optimization

Customer		
Cid	Ano	C_name
101	A1	Ram
102	A2	Harsh
103	A3	Deepak
104	A4	Gopal

Account	
Ano	Balance
A1	3000
A2	1000
A3	2000
A4	4000

Efficient Plan

$\Pi_{\text{Cust_Name}}(\underbrace{\sigma_{\text{Balance} < 2500}(\text{account})}_{2 \text{ records}} \bowtie \underbrace{\text{customer}}_{4 \text{ records}})$

$\Pi_{\text{Cust_Name}}(\underbrace{\sigma_{\text{Balance} < 2500}(\text{Account})}_{4 \text{ records}} \bowtie \underbrace{\text{Customer}}_{4 \text{ records}})$

Transformation of relational Expression

The background of the slide features abstract, overlapping geometric shapes in various shades of green, ranging from light lime to dark forest green. These shapes are primarily located on the right side and bottom of the slide, creating a modern, layered effect. A thin, light gray line also extends diagonally across the lower right portion of the image.

Cascade of selection

Combined selection operation can be divided into sequence of individual selection.

Customer			
<u>Cid</u>	<u>Ano</u>	<u>Cust_name</u>	Balance
C01	1	Raj	3000
C02	2	Meet	1000
C03	3	Harsh	2000
C04	4	<u>Punit</u>	4000

Output			
<u>Cid</u>	<u>Ano</u>	<u>Cust_name</u>	Balance
C02	2	Meet	1000

$\sigma_{\text{Ano} < 3 \wedge \text{Balance} < 2000}(\text{Customer})$

=

$\sigma_{\text{Ano} < 3}(\sigma_{\text{Balance} < 2000}(\text{Customer}))$

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

Selection operation

Selection operation are commutative

Customer			
<u>Cid</u>	<u>Ano</u>	<u>Cust_name</u>	Balance
C01	1	Raj	3000
C02	2	Meet	1000
C03	3	Harsh	2000
C04	4	<u>Punit</u>	4000

Output			
<u>Cid</u>	<u>Ano</u>	<u>Cust_name</u>	Balance
C02	2	Meet	1000

$\sigma_{\text{Ano} < 3}(\sigma_{\text{Balance} < 2000}(\text{Customer}))$

=

$\sigma_{\text{Balance} < 2000}(\sigma_{\text{Ano} < 3}(\text{Customer}))$

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

Project operation

- If more than one projection operation is used in expression then only the outer projection operation is required.

Customer			
<u>Cid</u>	<u>Ano</u>	<u>Cust_name</u>	Balance
C01	1	Raj	3000
C02	2	Meet	1000
C03	3	Harsh	2000
C04	4	<u>Punit</u>	4000

Output	
<u>Cust_name</u>	
Raj	
Meet	
Harsh	
<u>Punit</u>	

$$\Pi_{\text{Cust_name}} (\Pi_{\text{Ano, Cust_name}} (\text{Customer})) = \Pi_{\text{Cust_name}} (\text{Customer})$$

$$\Pi_{L1} (\Pi_{L2} (\dots (\Pi_{Ln} (E)) \dots)) = \Pi_{L1} (E)$$

Join

- ▶ Natural join operations are associative

$$(E1 \bowtie E2) \bowtie E3 = E1 \bowtie (E2 \bowtie E3)$$

Union and Intersection

- ▶ Set operations union and intersection are commutative.
- ▶ Set operation union and intersection are associative.

Customer	Employee	Student	Union	Intersect
Cust_name	Emp_name	Emp_name	Output	Output
Raj	Meet	Raj	Raj	Meet
Meet	Suresh	Meet	Meet	
			Suresh	

Example

Consider the Relational Algebra expression given below:

$(X \bowtie_{\theta} Y) \cap (X \bowtie_{\theta} Z)$ where X, Y, and Z are relational algebra expressions.

Identify the correct equivalent Relational Algebra expression.

- a) $X \bowtie_{\theta} (Y \bowtie_{\theta} Z)$
- b) $(X \bowtie_{\theta} Y) - Z$
- c) $(X \bowtie_{\theta} Y) - (X \bowtie_{\theta} Z)$
- d) $X \bowtie_{\theta} (Y \cap Z)$

Join Operation

- ▶ Several different algorithms to implement joins
 - ▶ Nested-loop join
 - ▶ Block nested-loop join
 - ▶ Indexed nested-loop join
 - ▶ Merge-join
 - ▶ Hash-join
- ▶ Choice based on cost estimate

Nested-Loop Join

- ▶ To compute the theta join $r \bowtie_{\theta} s$
for each tuple t_r in r do begin
for each tuple t_s in s do begin
test pair (t_r, t_s) to see if they satisfy the join condition θ
if they do, add $t_r \cdot t_s$ to the result.
end
end
- ▶ r is called the **outer relation** and s the **inner relation** of the join.
- ▶ Requires no indices and can be used with any kind of join condition.
- ▶ Expensive since it examines every pair of tuples in the two relations.

Nested-Loop Join (Cont.)

- ▶ In the worst case, if there is enough memory only to hold one block of each relation, the estimated cost is
$$\begin{array}{ll} n_r * b_s + b_r & \text{block transfers, plus} \\ n_r + b_r & \text{seeks} \end{array}$$
- ▶ N_r is the number of records in relation r , b_r is the number of blocks in which they exist
- ▶ If the smaller relation fits entirely in memory, use that as the inner relation.
 - ▶ Reduces cost to $b_r + b_s$ block transfers and 2 seeks
- ▶ Examples use the following information
 - ▶ Number of records of *student*: 5,000 *takes*: 10,000
 - ▶ Number of blocks of *student*: 100 *takes*: 400
- ▶ Assuming worst case memory availability cost estimate is
 - ▶ with *student* as outer relation:
 - ▶ $5000 * 400 + 100 = 2,000,100$ block transfers,
 - ▶ $5000 + 100 = 5100$ seeks
 - ▶ with *takes* as the outer relation
 - ▶ $10000 * 100 + 400 = 1,000,400$ block transfers and 10,400 seeks
- ▶ If smaller relation (*student*) fits entirely in memory, the cost estimate will be 500 block transfers.

Recap

- ▶ Query processing
- ▶ Measures of Query Cost
- ▶ Evaluation of expressions
- ▶ Query representation